

Performance-Aware Energy-Efficient GPU Frequency Selection using DNN-based Models

Artifact Description

The `environment.txt` file in `AD` directory includes the environment data of the NVIDIA GA100 node which is generated by machine using the bash script available at: <https://github.com/SC-Tech-Program/Author-Kit>.

Artifact Description

The data collection, deep learning-based model development, analyses and evaluations can be reproduced using the constructs available in `AE` directory. The description of each construct is as follows:

- **Data Collection:** allows acquisition of data for real applications (`LAMMPS`, `NAMD`, `GROMACS`, `LSTM`), `SPEC ACCEL`, `DGEMM`, and `STREAM`. The key components involved in data collection are `launch.sh`, `profile.py`, and `control.sh`.
- **Datasets:** A placeholder to include the aggregated data set for real applications, `SPEC ACCEL`, `DGEMM`, and `STREAM` (the dataset is removed due to size concerns).
- **Figures:** contains the diagrams produced in this study. All figures are reproducible using the data in `datasets` and `analysis.ipynb` notebook.
- **Models:** include the code to train, validate, and predict the power and performance models for different architectures.
- **Predictive analyses:** `analysis.ipynb` Jupyter notebook consists of the Python code which uses data in `datasets` and generates the plots available in `figures` and stores evaluated data in `results` directory.
- **Results:** contain numeric data generated during evaluation. These include optimal profiles (power, performance, frequency, energy, energy savings, performance degradation) for each application.